

DP-900.exam-labs.premium.exam.92q

Number: DP-900
Passing Score: 800
Time Limit: 120 min
File Version: 2.0

ExamLabs

DP-900

Microsoft Azure Data Fundamentals (beta)

Version 2.0

Sections

1. Describe core data concepts
2. Describe how to work with relational data on Azure
3. Describe how to work with non-relational data on Azure
4. Describe an analytics workload on Azure

Exam A

QUESTION 1 HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Descriptive analytics tells you

	▼
what is most likely to occur in the future.	
what occurred in the past.	
which actions you can perform to affect outcomes.	
why something occurred in the past.	

Correct Answer:

Answer Area

Descriptive analytics tells you

	▼
what is most likely to occur in the future.	
what occurred in the past.	
which actions you can perform to affect outcomes.	
why something occurred in the past.	

Section: Describe core data concepts
Explanation

Explanation/Reference:

Reference:

<https://demand-planning.com/2020/01/20/the-differences-between-descriptive-diagnostic-predictive-cognitive-analytics/>

QUESTION 2 HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input type="radio"/>
Normalizing a database reduces data redundancy.	<input type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input checked="" type="radio"/>
Normalizing a database reduces data redundancy.	<input checked="" type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe core data concepts
Explanation

Explanation/Reference:

Reference:

<https://www.sqlshack.com/what-is-database-normalization-in-sql-server/>

QUESTION 3

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

An extract, transform, and load (ETL) process requires

- | | |
|--|---|
| | ▼ |
| a matching schema in the data source and the data target. | |
| a target data store powerful enough to transform data. | |
| data that is fully processed before being loaded to the target data store. | |
| that the data target be a relational database. | |

Correct Answer:

Answer Area

An extract, transform, and load (ETL) process requires

- | | |
|--|---|
| | ▼ |
| a matching schema in the data source and the data target. | |
| a target data store powerful enough to transform data. | |
| data that is fully processed before being loaded to the target data store. | |
| that the data target be a relational database. | |

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

In the ELT pipeline, the transformation occurs in the target data store. ELT only works well when the target system is powerful enough to transform the data efficiently.

Incorrect Answers:

- The data does not need to be fully processed: Often, the three ETL phases are run in parallel to save time. For example, while data is being extracted, a transformation process could be working on data already received and prepare it for loading, and a loading process can begin working on the prepared data, rather than waiting for the entire extraction process to complete.
- The target does need to be a relational database.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl>

QUESTION 4

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

In batch processing,

	▼
data is always inserted one row at a time.	
data is processed in real-time.	
latency is expected.	
processing can only execute serially.	

Correct Answer:

Answer Area

In batch processing,

	▼
data is always inserted one row at a time.	
data is processed in real-time.	
latency is expected.	
processing can only execute serially.	

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

<https://www.bmc.com/blogs/what-is-batch-processing-batch-processing-explained/>

QUESTION 5

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Transcribing audio files is an example of

	▼
cognitive	
descriptive	
predictive	
prescriptive	

analytics.

Correct Answer:

Answer Area

Transcribing audio files is an example of

	▼
cognitive	
descriptive	
predictive	
prescriptive	

 analytics.

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

<https://azure.microsoft.com/en-us/services/cognitive-services/speech-services/>

QUESTION 6

DRAG DROP

Match the types of analytics that can be used to answer the business questions.

To answer, drag the appropriate analytics type from the column on the left to its question on the right. Each analytics type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Analytics Types

Cognitive
Diagnostic
Descriptive
Predictive
Prescriptive

Answer Area

	Why did sales increase last month?
	How do I allocate my budget to buy different inventory items?
	Which people are mentioned in a company's business documents?

Correct Answer:

Analytics Types

Cognitive
Diagnostic
Descriptive
Predictive
Prescriptive

Answer Area

Diagnostic	Why did sales increase last month?
Prescriptive	How do I allocate my budget to buy different inventory items?
Descriptive	Which people are mentioned in a company's business documents?

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

Box 1: Diagnostic

Diagnostic Analytics: At this stage you can begin to answer some of those why questions. Historical data can begin to be measured against other data to answer the question of why something happened in the past. This is the process of gathering and interpreting different data sets to identify anomalies, detect patterns, and determine relationships.

Box 2: Prescriptive

Prescriptive analytics is a combination of data, mathematical models, and various business rules to infer actions to influence future desired outcomes.

Incorrect Answer:

Predictive analytics, broadly speaking, is a category of business intelligence that uses descriptive and predictive variables from the past to analyze and identify the likelihood of an unknown future outcome

Box 3: Descriptive

Generally speaking, data analytics comes in four types:

Descriptive, to answer the question: What's happening?

Diagnostic, to answer the question: Why's happening?

Predictive, to answer the question: What will happen?

Prescriptive, to answer the question: What actions should we take?

**Reference:**

<https://demand-planning.com/2020/01/20/the-differences-between-descriptive-diagnostic-predictive-cognitive-analytics/>

<https://azure.microsoft.com/en-us/blog/answering-whats-happening-whys-happening-and-what-will-happen-with-iot-analytics/>

QUESTION 7**HOTSPOT**

You have the following JSON document.


```
"customer" : {
  "first name" : "Ben",
  "last name" : "Smith",
  "address" : {
    "line 1" : "161 Azure Ln",
    "line 2" : "Palo Alto",
    "ZIP code" : "54762"
  },
  "social media": [
    {
      "service" : "twitter",
      "handle" : "@bensmith"
    },
    {
      "service" : "linkedin",
      "handle" : "bensmith"
    }
  ],
  "phone numbers": [
    {
      "type" : "mobile",
      "number" : "555-555-555"
    }
  ]
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the JSON document.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Customer is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Address is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Social media is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Correct Answer:

Answer Area

Customer is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Address is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Social media is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Section: Describe core data concepts
Explanation

Explanation/Reference:

Reference:

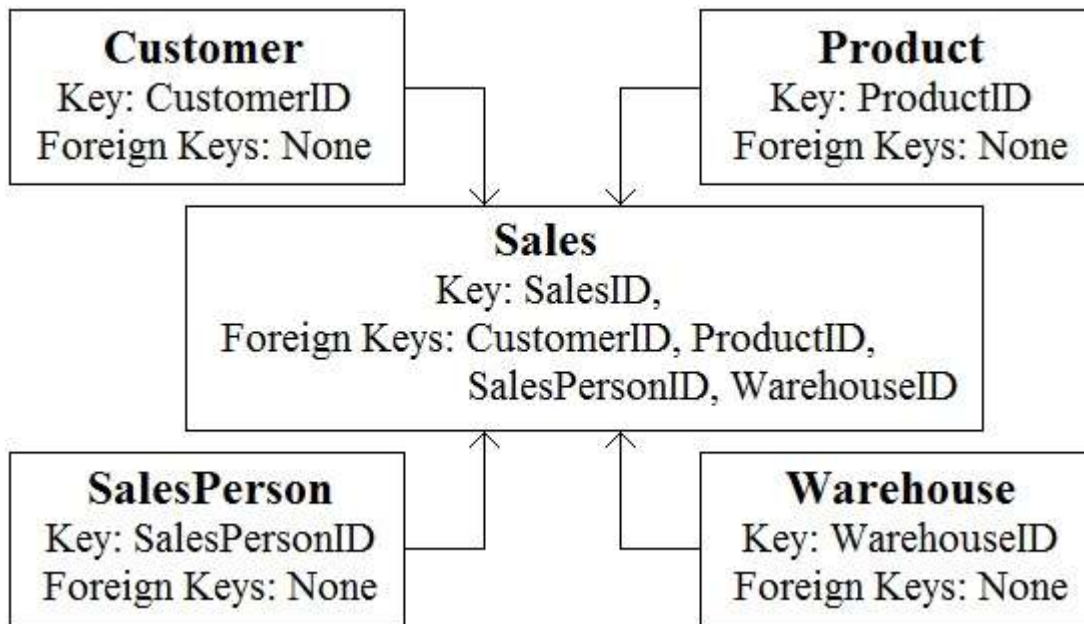
https://www.w3schools.com/js/js_json_arrays.asp

https://www.w3schools.com/js/js_json_objects.asp

QUESTION 8

HOTSPOT

You are reviewing the data model shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point

Hot Area:

Answer Area

The data model is a [answer choice].

	▼
transactional model	
star schema	
snowflake schema	

Customer is a [answer choice] table.

	▼
fact	
dimension	
bridge	

Correct Answer:

Answer Area

The data model is a [answer choice].

▼
transactional model
star schema
snowflake schema

Customer is a [answer choice] table.

▼
fact
dimension
bridge

Section: Describe core data concepts

Explanation


Explanation/Reference:

Explanation:

Box 1: star schema

In computing, the star schema is the simplest style of data mart schema and is the approach most widely used to develop data warehouses and dimensional data marts. The star schema consists of one or more fact tables referencing any number of dimension tables. The star schema is an important special case of the snowflake schema, and is more effective for handling simpler queries.

Example:




Incorrect Answers:

The data in the question is not normalized.

The snowflake schema is a variation of the star schema, featuring normalization of dimension tables.

Example:



Note: A snowflake schema is a logical arrangement of tables in a multidimensional database such that the entity relationship diagram resembles a snowflake shape. The snowflake schema is represented by centralized fact tables which are connected to multiple dimensions.[citation needed]. "Snowflaking" is a method of normalizing the dimension tables in a star schema. When it is completely normalized along all the dimension tables, the resultant structure resembles a snowflake with the fact table in the middle.

Box 2: dimension

The star schema consists of one or more fact tables referencing any number of dimension tables.

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-tables-overview>

https://en.wikipedia.org/wiki/Star_schema

https://en.wikipedia.org/wiki/Snowflake_schema

<https://azure.microsoft.com/en-us/blog/data-models-within-azure-analysis-services-and-power-bi/>

QUESTION 9

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

The massively parallel processing (MPP) engine of Azure Synapse Analytics

	▼
distributes processing across compute nodes.	
distributes processing across control nodes.	
redirects client connections across compute nodes.	
redirects client connections across control nodes.	

Correct Answer:

Answer Area

The massively parallel processing (MPP) engine of Azure Synapse Analytics

	▼
distributes processing across compute nodes.	
distributes processing across control nodes.	
redirects client connections across compute nodes.	
redirects client connections across control nodes.	

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/massively-parallel-processing-mpp-architecture>

QUESTION 10

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

	▼
A clustered index	
A FileTable	
A foreign key	
A stored procedure	

is an object associated with a table that sorts and stores the data rows in the table based on their key values.

Correct Answer:

Answer Area

	▼
A clustered index	
A FileTable	
A foreign key	
A stored procedure	

is an object associated with a table that sorts and stores the data rows in the table based on their key values.

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/clustered-and-nonclustered-indexes-described?view=sql-server-ver15>

QUESTION 11

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

A relational database is appropriate for scenarios that involve a high volume of

	▼
changes to relationships between entities	
geographically distributed writes	
transactional writes	
writes that have varying data structures	

Correct Answer:

Answer Area

A relational database is appropriate for scenarios that involve a high volume of

	▼
changes to relationships between entities	
geographically distributed writes	
transactional writes	
writes that have varying data structures	

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

Disadvantages of non-relational databases include: Data Consistency — non-relational databases do not perform ACID transactions.

Note: Relational databases are optimized for writes. They are optimized for consistency and availability. Advantages of relational databases include simplicity, ease of data retrieval, data integrity, and flexibility.

Incorrect Answers:

Use a relational database when data that you work with is structured, and the structure is not subject to frequent changes.

Use Cloud storage (no relational database) for geographically distributed writes.

Reference:

<https://towardsdatascience.com/choosing-the-right-database-c45cd3a28f77>

QUESTION 12

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Batch processing can output data to a file store	<input type="radio"/>	<input type="radio"/>
Batch processing can output data to a relational database	<input type="radio"/>	<input type="radio"/>
Batch processing can output data to a NoSQL database	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Batch processing can output data to a file store	<input checked="" type="radio"/>	<input type="radio"/>
Batch processing can output data to a relational database	<input type="radio"/>	<input checked="" type="radio"/>
Batch processing can output data to a NoSQL database	<input type="radio"/>	<input checked="" type="radio"/>

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

Box 1: Yes

Big data solutions often use long-running batch jobs to filter, aggregate, and otherwise prepare the data for analysis. Usually these jobs involve reading source files from scalable storage (like HDFS, Azure Data Lake Store, and Azure Storage), processing them, and writing the output to new files in scalable storage.

Box 2: No

Box 3: No

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/big-data/batch-processing>

QUESTION 13

DRAG DROP

Your company plans to load data from a customer relationship management (CRM) system to a data warehouse by using an extract, load, and transform (ELT) process.

Where does data processing occur for each stage of the ELT process? To answer, drag the appropriate locations to the correct stages. Each location may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

NOTE: Each correct selection is worth one point.

Select and Place:

Locations

An in-memory data integration tool

The CRM system

The data warehouse

Answer Area

Extract:

Location

Load:

Location

Transform:

Location

Correct Answer:

Locations

An in-memory data integration tool

The CRM system

The data warehouse

Answer Area

Extract:

The CRM system

Load:

The data warehouse

Transform:

An in-memory data integration tool

Section: Describe core data concepts

Explanation

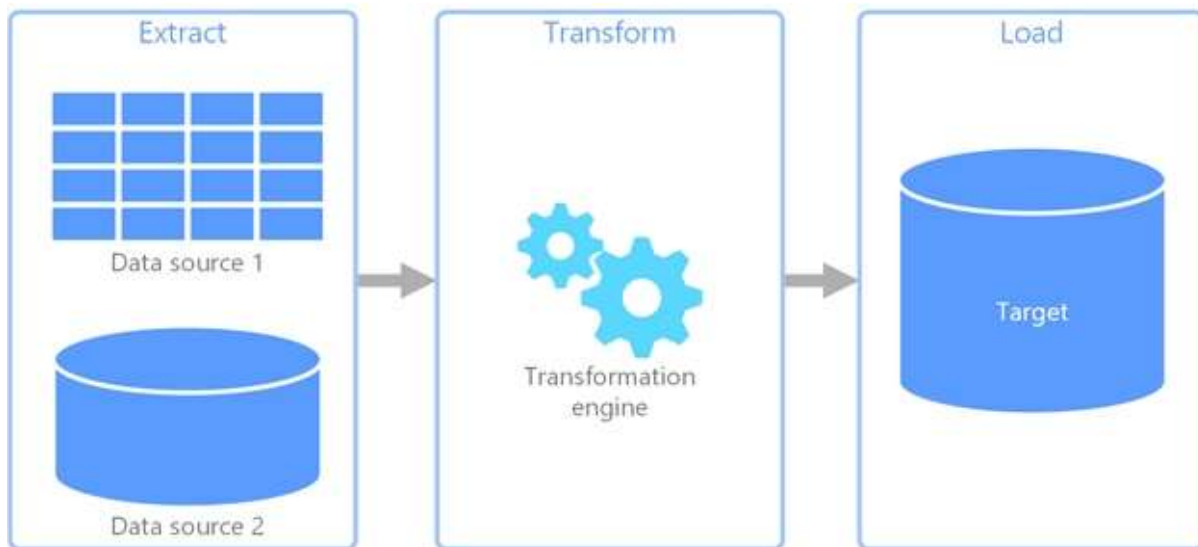
Explanation/Reference:

Explanation:

Box 1: The CRM system
Data is extracted from the CRM system.

Box 2: The data warehouse
Data is loaded to the data warehouse.

Box 3: An in-memory data integration tool
The data transformation that takes place usually involves various operations, such as filtering, sorting, aggregating, joining data, cleaning data, deduplicating, and validating data.



Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl>

QUESTION 14

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

A visualization that shows a university's current student enrollment versus the maximum capacity is an example of

	▼	analytics.
cognitive		
descriptive		
predictive		
prescriptive		

Correct Answer:

Answer Area

A visualization that shows a university's current student enrollment versus the maximum capacity is an example of

	▼	analytics.
cognitive		
descriptive		
predictive		
prescriptive		

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

Generally speaking, data analytics comes in four types (Figure 1):

Descriptive, to answer the question: What's happening?

Diagnostic, to answer the question: Why's happening?

Predictive, to answer the question: What will happen?

Prescriptive, to answer the question: What actions should we take?

Reference:

<https://azure.microsoft.com/en-us/blog/answering-whats-happening-whys-happening-and-what-will-happen-with-iot-analytics/>

QUESTION 15

DRAG DROP

Match the types of visualizations to the appropriate descriptions.

To answer, drag the appropriate visualization type from the column on the left to its description on the right. Each visualization type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Visualizations	Answer Area
Key influencer	Visualization A chart of colored, nested rectangles that displays individual data points represented by the size and color of a relative rectangle.
Scatter	Visualization A chart that displays the major contributors of a selected result or value.
Treemap	Visualization A chart that shows the relationship between two numerical values.

Correct Answer:

Visualizations	Answer Area
Key influencer	Treemap A chart of colored, nested rectangles that displays individual data points represented by the size and color of a relative rectangle.
Scatter	Key influencer A chart that displays the major contributors of a selected result or value.
Treemap	Scatter A chart that shows the relationship between two numerical values.

Section: Describe core data concepts


Explanation

Explanation/Reference:

Explanation:

Box 1: Tree map

Treemaps are charts of colored rectangles, with size representing value. They can be hierarchical, with rectangles nested within the main rectangles.



Box 2: Key influencer

A key influencer chart displays the major contributors to a selected result or value.

Box 3: Scatter

Scatter and Bubble charts display relationships between 2 (scatter) or 3 (bubble) quantitative measures -- whether or not, in which order, etc.

QUESTION 16

You need to create an Azure Storage account.

Data in the account must replicate outside the Azure region automatically.

Which two types of replication can you use for the storage account? Each correct answer presents a complete solution.

NOTE: Each correct selection is worth one point.

- A. zone-redundant storage (ZRS)
- B. read-access geo-redundant storage (RA-GRS)
- C. locally-redundant storage (LRS)
- D. geo-redundant storage (GRS)

Correct Answer: BD

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

D: Azure Storage offers two options for copying your data to a secondary region:

- Geo-redundant storage (GRS)
- Geo-zone-redundant storage (GZRS)

B: With GRS or GZRS, the data in the secondary region isn't available for read or write access unless there is a failover to the secondary region. For read access to the secondary region, configure your storage account to use read-access geo-redundant storage (RA-GRS) or read-access geo-zone-redundant storage (RA-GZRS).

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy#redundancy-in-a-secondary-region>

QUESTION 17

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings.	<input type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide administrators with the ability to control and update the operating system version.	<input type="radio"/>	<input type="radio"/>
All platform as a service (PaaS) database offerings in Azure can be paused to reduce costs.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure require less setup and configuration effort than infrastructure as a service (IaaS) database offerings.	<input checked="" type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide administrators with the ability to control and update the operating system version.	<input checked="" type="radio"/>	<input type="radio"/>
All platform as a service (PaaS) database offerings in Azure can be paused to reduce costs.	<input type="radio"/>	<input checked="" type="radio"/>

Section: Describe core data concepts Explanation

Explanation/Reference:

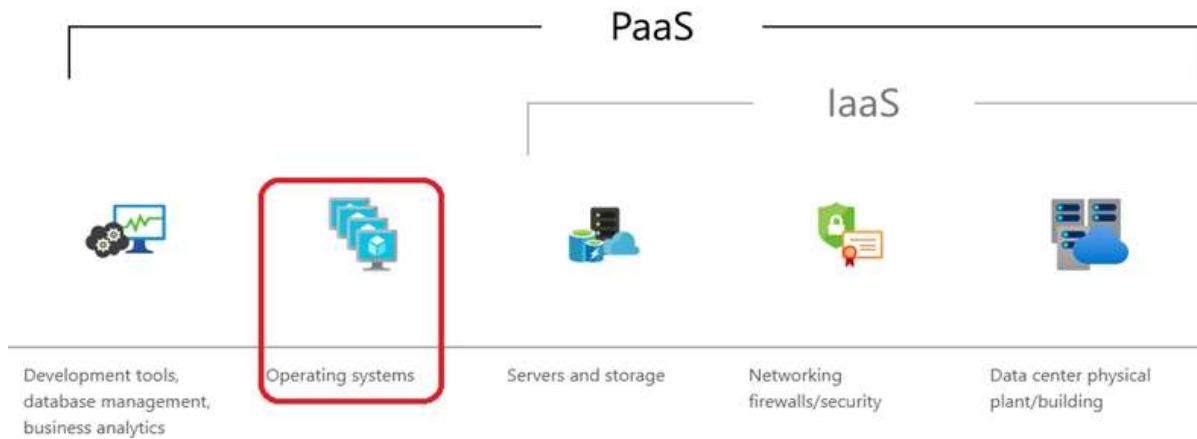
Explanation:

Box 1: Yes

Like IaaS, PaaS includes infrastructure – servers, storage, and networking – but also middleware, development tools, business intelligence (BI) services, database management systems, and more. PaaS is designed to support the complete web application lifecycle: building, testing, deploying, managing, and updating. PaaS allows you to avoid the expense and complexity of buying and managing software licenses, the underlying application infrastructure and middleware, container orchestrators such as Kubernetes, or the development tools and other resources

Box 2: Yes

You manage the applications and services you develop, and the cloud service provider typically manages everything else.



Box 3: No

There really is no way to pause / stop billing for your Azure SQL Database.

Microsoft's official answer "Yes, you can export your database. Delete the Azure SQL database and that will pause billing. Then when you need it you can create a new database and import your previously exported DB."

Reference:

<https://azure.microsoft.com/en-us/overview/what-is-paas>

<https://feedback.azure.com/forums/217321-sql-database/suggestions/6931152-please-add-ability-to-temporarily-turn-off-on-sql>

QUESTION 18

Which statement is an example of Data Manipulation Language (DML)?

- A. REVOKE
- B. DISABLE
- C. INSERT
- D. GRANT

Correct Answer: C

Section: Describe core data concepts

Explanation

Explanation/Reference:

Explanation:

Data Manipulation Language (DML) statements:

- DELETE
- INSERT
- UPDATE

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-reference-tsql-statements>

QUESTION 19

Which command-line tool can you use to query Azure SQL databases?

- A. sqlcmd
- B. bcp
- C. azdata
- D. Azure CLI

Correct Answer: A

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

The sqlcmd utility lets you enter Transact-SQL statements, system procedures, and script files at the command prompt.

Incorrect Answers:

B: The bulk copy program utility (bcp) bulk copies data between an instance of Microsoft SQL Server and a data file in a user-specified format.

D: The Azure CLI is the defacto tool for cross-platform and command-line tools for building and managing Azure resources.

Reference:

<https://docs.microsoft.com/en-us/sql/tools/overview-sql-tools?view=sql-server-ver15>

QUESTION 20

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure SQL Database includes a fully managed backup service.	<input type="radio"/>	<input type="radio"/>
Azure SQL Database has built-in high availability.	<input type="radio"/>	<input type="radio"/>
Azure SQL Database can use Azure Advanced Threat Protection (ATP).	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Azure SQL Database includes a fully managed backup service.	<input checked="" type="radio"/>	<input type="radio"/>
Azure SQL Database has built-in high availability.	<input checked="" type="radio"/>	<input type="radio"/>
Azure SQL Database can use Azure Advanced Threat Protection (ATP).	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/sql-database-paas-overview>

<https://azure.microsoft.com/en-us/blog/announcing-sql-atp-and-sql-vulnerability-assessment-general-availability/>

QUESTION 21

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	<input type="radio"/>	<input type="radio"/>
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	<input type="radio"/>	<input type="radio"/>
You can use MySQL Workbench to query Azure Database for MariaDB databases.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	<input checked="" type="radio"/>	<input type="radio"/>
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	<input checked="" type="radio"/>	<input type="radio"/>
You can use MySQL Workbench to query Azure Database for MariaDB databases.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/big-data-cluster/connect-to-big-data-cluster?view=sql-server-ver15>

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-query-ssms>

<https://docs.microsoft.com/en-us/azure/mariadb/connect-workbench>

QUESTION 22
HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure provide built-in high availability.	<input type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide configurable scaling options.	<input type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure provide built-in high availability.	<input checked="" type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide configurable scaling options.	<input checked="" type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/sql-database-paas-overview>

QUESTION 23
HOTSPOT

You have the following SQL query.

```
INSERT INTO dbo.Products (ProductID, ProductName, Price, ProductDescription)
VALUES (1, 'Clamp', 12.48, 'Workbench clamp') ;
```

What are `dbo.Products` and `ProductName`? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

<code>dbo.Products</code> :	<div><div></div><div>▼</div></div> <div>A column</div> <div>A database</div> <div>A table</div> <div>An index</div>
<code>ProductName</code> :	<div><div></div><div>▼</div></div> <div>A column</div> <div>A database</div> <div>A table</div> <div>An index</div>

Correct Answer:

Answer Area

<code>dbo.Products</code> :	<div><div></div><div>▼</div></div> <div>A column</div> <div>A database</div> <div>A table</div> <div>An index</div>
<code>ProductName</code> :	<div><div></div><div>▼</div></div> <div>A column</div> <div>A database</div> <div>A table</div> <div>An index</div>

Section: Describe how to work with relational data on Azure
Explanation

Explanation/Reference:

QUESTION 24
HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	<input type="radio"/>	<input type="radio"/>
You need a Microsoft 365 subscription to create an Azure SQL database.	<input type="radio"/>	<input type="radio"/>
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	<input type="radio"/>	<input checked="" type="radio"/>
You need a Microsoft 365 subscription to create an Azure SQL database.	<input type="radio"/>	<input checked="" type="radio"/>
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with relational data on Azure
Explanation

Explanation/Reference:

Reference:

<https://azure.microsoft.com/en-gb/blog/hot-patching-sql-server-engine-in-azure-sql-database/>

<https://azure.microsoft.com/en-us/services/sql-database/#product-overview>

QUESTION 25

Which statement is an example of Data Definition Language (DDL)?

- A. SELECT
- B. JOIN
- C. MERGE
- D. CREATE

Correct Answer: D

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://www.geeksforgeeks.org/difference-between-ddl-and-dml-in-dbms/>

QUESTION 26

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Transparent Data Encryption (TDE) encrypts

	▼
a column to protect data at rest and in transit.	
queries and their results in order to protect data in transit.	
the database to protect data at rest.	
the server to protect data at rest.	

Correct Answer:

Answer Area

Transparent Data Encryption (TDE) encrypts

	▼
a column to protect data at rest and in transit.	
queries and their results in order to protect data in transit.	
the database to protect data at rest.	
the server to protect data at rest.	

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/transparent-data-encryption-tde-overview?tabs=azure-portal>

QUESTION 27

You need to ensure that users use multi-factor authentication (MFA) when connecting to an Azure SQL database.

Which type of authentication should you use?

- A. service principal authentication
- B. Azure Active Directory (Azure AD) authentication
- C. SQL authentication
- D. certificate authentication

Correct Answer: B

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/authentication-mfa-ssms-overview>

QUESTION 28

What is a benefit of hosting a database on Azure SQL managed instance as compared to an Azure SQL database?

- A. built-in high availability
- B. native support for cross-database queries and transactions
- C. system-initiated automatic backups
- D. support for encryption at rest

Correct Answer: B

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/features-comparison>

QUESTION 29

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

By default, each Azure SQL database is protected by

	▼
a network security group (NSG).	
a server-level firewall.	
Azure Firewall.	
Azure Front Door.	

Correct Answer:

Answer Area

By default, each Azure SQL database is protected by

	▼
a network security group (NSG).	
a server-level firewall.	
Azure Firewall.	
Azure Front Door.	

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

When you create a new server in Azure SQL Database or Azure Synapse Analytics named mysqlserver, for example, a server-level firewall blocks all access to the public endpoint for the server

Reference:

<https://docs.microsoft.com/en-us/azure/security/fundamentals/infrastructure-sql>

QUESTION 30

You need to design and model a database by using a graphical tool that supports project-oriented offline database development.

What should you use?

- A. Microsoft SQL Server Data Tools (SSDT)
- B. Microsoft SQL Server Management Studio (SSMS)
- C. Azure Databricks
- D. Azure Data Studio

Correct Answer: A

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/ssdt/project-oriented-offline-database-development?view=sql-server-ver15>

QUESTION 31

DRAG DROP

Match the security components to the appropriate scenarios.

To answer, drag the appropriate component from the column on the left to its scenario on the right. Each component may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Components

Authentication
Firewall
Encryption

Answer Area

Prevent access to an Azure SQL database from another network.
 Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.
 Ensure that sensitive data never appears as plain text in an Azure SQL database.

Correct Answer:

Components

Authentication
Firewall
Encryption

Answer Area

Firewall Prevent access to an Azure SQL database from another network.
Authentication Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.
Encryption Ensure that sensitive data never appears as plain text in an Azure SQL database.

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/database-engine/configure-windows/configure-a-windows-firewall-for->

[database-engine-access?view=sql-server-ver15](#)

<https://docs.microsoft.com/en-us/azure/azure-sql/database/authentication-aad-overview>

<https://docs.microsoft.com/en-us/azure/azure-sql/database/always-encrypted-certificate-store-configure>

QUESTION 32

You have a transactional application that stores data in an Azure SQL managed instance.

When should you implement a read-only database replica?

- A. You need to generate reports without affecting the transactional workload.
- B. You need to audit the transactional application.
- C. You need to implement high availability in the event of a regional outage.
- D. You need to improve the recovery point objective (RPO).

Correct Answer: A

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Use read-only replicas to offload read-only query workloads.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/read-scale-out>

QUESTION 33

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

A relational database must be used when

	▼
a dynamic schema is required.	
data will be stored as key/value pairs.	
storing large images and videos.	
strong consistency guarantees are required.	

Correct Answer:

Answer Area

A relational database must be used when

	▼
a dynamic schema is required.	
data will be stored as key/value pairs.	
storing large images and videos.	
strong consistency guarantees are required.	

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

QUESTION 34

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Relational data uses

	▼
--	---

 to enforce relationships between different tables.

collections
columns
keys
partitions

Correct Answer:

Answer Area

Relational data uses

	▼
--	---

 to enforce relationships between different tables.

collections
columns
keys
partitions

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://teachcomputerscience.com/relational-databases/>

QUESTION 35

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

	▼
--	---

 is a virtual table that contains content defined by a query.

A heap
A stored procedure
A view
An index

Correct Answer:

Answer Area

	▼	is a virtual table that contains content defined by a query.
A heap		
A stored procedure		
A view		
An index		

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/views/views>

QUESTION 36

You need to query a table named Products in an Azure SQL database.

Which three requirements must be met to query the table from the internet? Each correct answer presents part of the solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. You must be assigned the Reader role for the resource group that contains the database.
- B. You must have `SELECT` access to the Products table.
- C. You must have a user in the database.
- D. You must be assigned the Contributor role for the resource group that contains the database.
- E. Your IP address must be allowed to connect to the database.

Correct Answer: BCE

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Incorrect Answers:

A, D: Resource group permissions is not required to query an Azure SQL database table.

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/security/authentication-access/getting-started-with-database-engine-permissions?view=sql-server-ver15>

QUESTION 37

You have an inventory management database that contains the following table.

ProductName	Quantity
Product1	100
Product2	129
Product3	176

Which statement should you use in a SQL query to change the inventory quantity of Product1 to 270?

- A. INSERT
- B. MERGE
- C. UPDATE
- D. CREATE

Correct Answer: C

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/t-sql/queries/update-transact-sql?view=sql-server-ver15>

QUESTION 38

Your company needs to implement a relational database in Azure. The solution must minimize ongoing maintenance.

Which Azure service should you use?

- A. Azure HDInsight
- B. Azure SQL Database
- C. Azure Cosmos DB
- D. SQL Server on Azure virtual machines

Correct Answer: B

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://azure.microsoft.com/en-us/services/sql-database/#features>

QUESTION 39

You are writing a set of SQL queries that administrators will use to troubleshoot an Azure SQL database.

You need to embed documents and query results into a SQL notebook.

What should you use?

- A. Microsoft SQL Server Management Studio (SSMS)
- B. Azure Data Studio
- C. Azure CLI
- D. Azure PowerShell

Correct Answer: B

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://www.mssqltips.com/sqlservertip/5997/create-sql-server-notebooks-in-azure-data-studio/>

QUESTION 40

DRAG DROP

Match the terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Terms

Index
View
Table

Answer Area

	A database object that holds data
	A database object whose content is defined by a query
	A database object that helps improve the speed of data retrieval

Correct Answer:

Terms

Index
View
Table

Answer Area

Table	A database object that holds data
View	A database object whose content is defined by a query
Index	A database object that helps improve the speed of data retrieval

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Reference:

[https://en.wikipedia.org/wiki/Table_\(database\)#:~:text=A%20table%20is%20a%20collection,table%20format%20within%20a%20database.&text=In%20relational%20databases%2C%20and%20flat,a%20row%20and%20column%20intersect](https://en.wikipedia.org/wiki/Table_(database)#:~:text=A%20table%20is%20a%20collection,table%20format%20within%20a%20database.&text=In%20relational%20databases%2C%20and%20flat,a%20row%20and%20column%20intersect)

[https://en.wikipedia.org/wiki/View_\(SQL\)](https://en.wikipedia.org/wiki/View_(SQL))

https://en.wikipedia.org/wiki/Database_index#:~:text=A%20database%20index%20is%20a,maintain%20the%20index%20data%20structure

QUESTION 41

You have an e-commerce application that reads and writes data to an Azure SQL database.

Which type of processing does the application use?

- A. stream processing
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. Online Transaction Processing (OLTP)

Correct Answer: D

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

OLTP is designed to serve as a persistent data store for business or front-end applications. OLTP administers day to day transaction of an organization.

Reference:

<https://sqlwizard.blog/2020/03/15/sql-server-oltp-vs-olap/>

QUESTION 42

When can you use an Azure Resource Manager template?

- A. to automate the creation of an interdependent group of Azure resources in a repeatable way
- B. to apply Azure policies for multi-tenant deployments
- C. to provision Azure subscriptions
- D. to control which services and feature administrators and developers can deploy from the Azure portal

Correct Answer: A

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

You can automate deployments and use the practice of infrastructure as code. In code, you define the infrastructure that needs to be deployed

To implement infrastructure as code for your Azure solutions, use Azure Resource Manager templates (ARM templates). The template is a JavaScript Object Notation (JSON) file that defines the infrastructure and configuration for your project. The template uses declarative syntax, which lets you state what you intend to deploy without having to write the sequence of programming commands to create it. In the template, you specify the resources to deploy and the properties for those resources.

Reference:

<https://docs.microsoft.com/en-us/azure/azure-resource-manager/templates/overview>

QUESTION 43

You have an Azure SQL database that you access directly from the Internet.

You recently changed your external IP address.

After changing the IP address, you can no longer access the database. You can connect to other resources in Azure.

What is a possible cause of the issue?

- A. role-based access control (RBAC)
- B. Dynamic Host Configuration Protocol (DHCP)
- C. Domain Name Service (DNS)
- D. a database-level firewall

Correct Answer: D

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

The Azure SQL Database firewall lets you decide which IP addresses may or may not have access to either your Azure SQL Server or your Azure SQL database.

When creating an Azure SQL Database, the firewall needs to be configured before anyone will be able to access the database. By default, no external access to your SQL Database will be allowed until you explicitly assign permission by creating a firewall rule.

Reference:

<https://www.sqlshack.com/configuring-the-azure-sql-database-firewall/>

QUESTION 44

DRAG DROP

Match the tools to the appropriate descriptions.

To answer, drag the appropriate tool from the column on the left to its description on the right. Each tool may be used once, more than once, or not at all.

Select and Place:

Tools	Answer Area	
Azure Data Studio	Tool	A graphical tool for managing SQL Server or Azure SQL databases that supports access, configuration, management, and administration tasks.
Microsoft SQL Server Data Tools (SSDT)	Tool	A lightweight source code editor with an mssql extension that supports connections to SQL Server and a rich editing experience for T-SQL.
Microsoft SQL Server Management Studio (SSMS)	Tool	A lightweight editor that can run on-demand SQL queries and view and save results as text, JSON, or Microsoft Excel files.
Microsoft Visual Studio Code	Tool	A development tool for building Azure SQL databases, Microsoft SQL Server relational databases, SQL Server Analysis Services (SSAS) data models, SQL Server Integration Services (SSIS) packages, and SQL Server Reporting Services (SSRS) reports.

Correct Answer:

Tools	Answer Area	
Azure Data Studio	Microsoft SQL Server Management Studio (SSMS)	A graphical tool for managing SQL Server or Azure SQL databases that supports access, configuration, management, and administration tasks.
Microsoft SQL Server Data Tools (SSDT)	Microsoft Visual Studio Code	A lightweight source code editor with an mssql extension that supports connections to SQL Server and a rich editing experience for T-SQL.
Microsoft SQL Server Management Studio (SSMS)	Azure Data Studio	A lightweight editor that can run on-demand SQL queries and view and save results as text, JSON, or Microsoft Excel files.
Microsoft Visual Studio Code	Microsoft SQL Server Data Tools (SSDT)	A development tool for building Azure SQL databases, Microsoft SQL Server relational databases, SQL Server Analysis Services (SSAS) data models, SQL Server Integration Services (SSIS) packages, and SQL Server Reporting Services (SSRS) reports.

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: Microsoft SQL Server Management Studio (SSMS)

SQL Server Management Studio (SSMS) is an integrated environment for managing any SQL infrastructure, from SQL Server to Azure SQL Database.

Box 2: Microsoft Visual Studio Code

Visual Studio Code is a streamlined code editor with support for development operations like debugging, task running, and version control. It aims to provide just the tools a developer needs for a quick code-build-debug cycle and leaves more complex workflows to fuller featured IDEs, such as Visual Studio IDE.

Box 3: Azure Data Studio

Azure Data Studio offers a modern, keyboard-focused SQL coding experience that makes your everyday tasks easier with built-in features, such as multiple tab windows, a rich SQL editor, IntelliSense, keyword completion, code snippets, code navigation, and source control integration (Git). Run on-demand SQL queries, view and

save results as text, JSON, or Excel. Edit data, organize your favorite database connections, and browse database objects in a familiar object browsing experience.

Box 4: Microsoft SQL Server Data Tools (SSDT)

SQL Server Data Tools (SSDT) is a modern development tool for building SQL Server relational databases, databases in Azure SQL, Analysis Services (AS) data models, Integration Services (IS) packages, and Reporting Services (RS) reports. With SSDT, you can design and deploy any SQL Server content type with the same ease as you would develop an application in Visual Studio.

Reference:

<https://docs.microsoft.com/en-us/sql/ssms/download-sql-server-management-studio-ssms>

<https://code.visualstudio.com/docs/supporting/FAQ>

<https://docs.microsoft.com/en-us/sql/azure-data-studio/what-is-azure-data-studio>

<https://docs.microsoft.com/en-us/sql/ssdt/download-sql-server-data-tools-ssdt>

QUESTION 45

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Relational database tables contain columns and rows	<input type="radio"/>	<input type="radio"/>
Indexes in a relational database describe the data types in a table	<input type="radio"/>	<input type="radio"/>
A database view is a virtual table whose content is defined by a query	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Relational database tables contain columns and rows	<input checked="" type="radio"/>	<input type="radio"/>
Indexes in a relational database describe the data types in a table	<input type="radio"/>	<input checked="" type="radio"/>
A database view is a virtual table whose content is defined by a query	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with relational data on Azure
Explanation

Explanation/Reference:

Explanation:

Box 1: Yes

Tables are database objects that contain all the data in a database. In tables, data is logically organized in a row-and-column format similar to a spreadsheet. Each row represents a unique record, and each column represents a field in the record.

Box 2: No

An index is an on-disk structure associated with a table or view that speeds retrieval of rows from the table or view.

Box 3: Yes

A view is a virtual table whose contents are defined by a query. Like a table, a view consists of a set of named columns and rows of data.

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/tables/tables>

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/clustered-and-nonclustered-indexes-described>

<https://docs.microsoft.com/en-us/sql/relational-databases/views/views?view=sql-server-ver15>

QUESTION 46

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
When ingesting data from Azure Data Lake Storage across Azure regions, you will incur costs for bandwidth.	<input type="radio"/>	<input type="radio"/>
You can use blob, table, and file storage in the same Azure Storage account.	<input type="radio"/>	<input type="radio"/>
You implement Azure Data Lake Storage by creating an Azure Storage account.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
When ingesting data from Azure Data Lake Storage across Azure regions, you will incur costs for bandwidth.	<input checked="" type="radio"/>	<input type="radio"/>
You can use blob, table, and file storage in the same Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>
You implement Azure Data Lake Storage by creating an Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/data-lake-store/data-lake-store-get-started-portal>

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview>

<https://azure.microsoft.com/en-us/pricing/details/bandwidth/>

QUESTION 47

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Correct Answer:

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/create-graph-gremlin-console>

QUESTION 48

At which two levels can you set the throughput for an Azure Cosmos DB account? Each correct answer presents a complete solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. database
- B. item
- C. container
- D. partition

Correct Answer: AC

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/set-throughput>

QUESTION 49

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input checked="" type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/high-availability>

QUESTION 50

DRAG DROP

Match the types of data stores to the appropriate scenarios.

To answer, drag the appropriate data store type from the column on the left to its scenario on the right. Each data store type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Data Store Types

Graph
Object
Key/value

Answer Area

<input type="text"/>	Application users and their default language
<input type="text"/>	Medical images and their associated metadata
<input type="text"/>	Employee data that shows the relationships between employees

Correct Answer:

Data Store Types

Graph
Object
Key/value

Answer Area

Key/value	Application users and their default language
Object	Medical images and their associated metadata
Graph	Employee data that shows the relationships between employees

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

QUESTION 51

You have an Azure Cosmos DB account that uses the Core (SQL) API.

Which two settings can you configure at the container level? Each correct answer presents a complete solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. the throughput
- B. the read region
- C. the partition key
- D. the API

Correct Answer: AC

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://www.sqlshack.com/start-your-journey-with-azure-cosmos-db/>

QUESTION 52

Your company is designing a data store that will contain student data. The data has the following format.

StudentNumber	StudentInformation
7634634	First name: Ben Last: Smith Preferred Name: Benjamin
7634634	First Name: Dominik Last Name: Paiha Email Address: dpaiha@contoso.com MCP ID: 931817
7634636	First Name: Reshma Last Name: Patel Phone number: 514-555-1101
7634637	First Name: Yun-Feng Last Name: Peng

Which type of data store should you use?

- A. graph
- B. key/value
- C. object
- D. columnar

Correct Answer: D

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

QUESTION 53

Which storage solution supports role-based access control (RBAC) at the file and folder level?

- A. Azure Disk Storage
- B. Azure Data Lake Storage
- C. Azure Blob storage
- D. Azure Queue storage

Correct Answer: B

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-access-control>

QUESTION 54

You need to store data in Azure Blob storage for seven years to meet your company's compliance requirements. The retrieval time of the data is unimportant. The solution must minimize storage costs.

Which storage tier should you use?

- A. Archive
- B. Hot
- C. Cool

Correct Answer: A

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

https://cloud.netapp.com/blog/azure-blob-storage-pricing-the-complete-guide-azure-cvo-blg#H1_4

QUESTION 55

Which type of non-relational data store supports a flexible schema, stores data as JSON files, and stores the all the data for an entity in the same document?

- A. document
- B. columnar
- C. graph
- D. time series

Correct Answer: A

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

QUESTION 56

DRAG DROP

Match the Azure Cosmos DB APIs to the appropriate data structures.

To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

APIs

Cassandra API

Gremlin API

MongoDB API

Table API

Answer Area

Graph data

JSON documents

Key/value data

Correct Answer:

APIs

Cassandra API

Gremlin API

MongoDB API

Table API

Answer Area

Gremlin API

Graph data

MongoDB API

JSON documents

Table API

Key/value data

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/faq>

QUESTION 57

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,

enable the hierarchical namespace.

set Account kind to BlobStorage.

set Performance to Premium.

set Replication to Read-access geo-redundant storage (RA-GRS).

Correct Answer:

Answer Area

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,

	▼
enable the hierarchical namespace.	
set Account kind to BlobStorage.	
set Performance to Premium.	
set Replication to Read-access geo-redundant storage (RA-GRS).	

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-namespace>

QUESTION 58

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

You can query a graph database in Azure Cosmos DB

	▼
as a JSON document by using a SQL-like language.	
as a partitioned row store by using Cassandra Query Language (CQL).	
as a partitioned row store by using Language-Integrated Query (LINQ).	
as nodes and edges by using the Gremlin language.	

Correct Answer:

Answer Area

You can query a graph database in Azure Cosmos DB

	▼
as a JSON document by using a SQL-like language.	
as a partitioned row store by using Cassandra Query Language (CQL).	
as a partitioned row store by using Language-Integrated Query (LINQ).	
as nodes and edges by using the Gremlin language.	

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://www.sqlshack.com/graph-database-implementation-with-azure-cosmos-db-using-the-api/>

QUESTION 59

You manage an application that stores data in a shared folder on a Windows server.

You need to move the shared folder to Azure Storage.

Which type of Azure Storage should you use?

- A. queue
- B. blob
- C. file

D. table

Correct Answer: C

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Azure file shares can be mounted concurrently by cloud or on-premises deployments of Windows, Linux, and macOS. Azure file shares can also be cached on Windows Servers with Azure File Sync for fast access near where the data is being used.

Reference:

<https://azure.microsoft.com/en-us/services/storage/files/>

QUESTION 60

Your company is designing a database that will contain session data for a website. The data will include notifications, personalization attributes, and products that are added to a shopping cart.

Which type of data store will provide the lowest latency to retrieve the data?

- A. key/value
- B. graph
- C. columnar
- D. document

Correct Answer: C

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/analytical-data-stores>

QUESTION 61

You have an application that runs on Windows and requires access to a mapped drive.

Which Azure service should you use?

- A. Azure Files
- B. Azure Blob storage
- C. Azure Cosmos DB
- D. Azure Table storage

Correct Answer: A

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Azure Files is Microsoft's easy-to-use cloud file system. Azure file shares can be seamlessly used in Windows and Windows Server.

To use an Azure file share with Windows, you must either mount it, which means assigning it a drive letter or mount point path, or access it via its UNC path.

Reference:

<https://docs.microsoft.com/en-us/azure/storage/files/storage-how-to-use-files-windows>

QUESTION 62
HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account.	<input type="radio"/>	<input type="radio"/>
Partition keys are used in Azure Cosmos DB to optimize queries.	<input type="radio"/>	<input type="radio"/>
Items contained in the same Azure Cosmos DB logical partition can have different partition keys.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
The Azure Cosmos DB API is configured separately for each database in an Azure Cosmos DB account.	<input type="radio"/>	<input checked="" type="radio"/>
Partition keys are used in Azure Cosmos DB to optimize queries.	<input checked="" type="radio"/>	<input type="radio"/>
Items contained in the same Azure Cosmos DB logical partition can have different partition keys.	<input type="radio"/>	<input checked="" type="radio"/>

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: No

The API determines the type of account to create. Azure Cosmos DB provides five APIs: Core (SQL) and MongoDB for document data, Gremlin for graph data, Azure Table, and Cassandra. Currently, you must create a separate account for each API.

Box 2: Yes

Azure Cosmos DB uses partitioning to scale individual containers in a database to meet the performance needs of your application. In partitioning, the items in a container are divided into distinct subsets called logical partitions. Logical partitions are formed based on the value of a partition key that is associated with each item in a container.

Box 3: No

Logical partitions are formed based on the value of a partition key that is associated with each item in a

container.

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/partitioning-overview>

QUESTION 63

Your company is designing an application that will write a high volume of JSON data and will have an application-defined schema.

Which type of data store should you use?

- A. columnar
- B. key/value
- C. document
- D. graph

Correct Answer: B

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

A key/value store associates each data value with a unique key.

An application can store arbitrary data as a set of values. Any schema information must be provided by the application. The key/value store simply retrieves or stores the value by key.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

QUESTION 64

You need to recommend a non-relational data store that is optimized for storing and retrieving text files, videos, audio streams, and virtual disk images. The data store must store data, some metadata, and a unique ID for each file.

Which type of data store should you recommend?

- A. key/value
- B. columnar
- C. object
- D. document

Correct Answer: C

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Object storage is optimized for storing and retrieving large binary objects (images, files, video and audio streams, large application data objects and documents, virtual machine disk images). Large data files are also popularly used in this model, for example, delimiter file (CSV), parquet, and ORC. Object stores can manage extremely large amounts of unstructured data.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

QUESTION 65

Your company is designing a data store for internet-connected temperature sensors.

The collected data will be used to analyze temperature trends.

Which type of data store should you use?

- A. relational
- B. time series
- C. graph
- D. columnar

Correct Answer: B

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Explanation:

Time series data is a set of values organized by time. Time series databases typically collect large amounts of data in real time from a large number of sources. Updates are rare, and deletes are often done as bulk operations. Although the records written to a time-series database are generally small, there are often a large number of records, and total data size can grow rapidly.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

QUESTION 66

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	<input type="radio"/>	<input type="radio"/>
Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input type="radio"/>
Azure Data Factory orchestrates data integration workflows.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	<input checked="" type="radio"/>	<input type="radio"/>
Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input checked="" type="radio"/>
Azure Data Factory orchestrates data integration workflows.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/databricks/scenarios/what-is-azure-databricks>

<https://docs.microsoft.com/en-us/azure/analysis-services/analysis-services-overview>

<https://docs.microsoft.com/en-us/azure/data-factory/introduction>

QUESTION 67

Which scenario is an example of a streaming workload?

- A. sending transactions that are older than a month to an archive
- B. sending transactions daily from point of sale (POS) devices
- C. sending telemetry data from edge devices
- D. sending cloud infrastructure metadata every 30 minutes

Correct Answer: C

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

QUESTION 68

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Batch workloads

	▼
process data in memory, row-by-row.	
collect and process data at most once a day.	
process data as new data is received in near real-time.	
collect data and then process the data when a condition is met.	

Correct Answer:

Answer Area

Batch workloads

- process data in memory, row-by-row.
- collect and process data at most once a day.
- process data as new data is received in near real-time.
- collect data and then process the data when a condition is met.

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

QUESTION 69

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	<input type="radio"/>	<input type="radio"/>
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	<input checked="" type="radio"/>	<input type="radio"/>
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input checked="" type="radio"/>	<input type="radio"/>
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input checked="" type="radio"/>

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

QUESTION 70

You need to gather real-time telemetry data from a mobile application.

Which type of workload describes this scenario?

- A. Online Transaction Processing (OLTP)
- B. batch
- C. massively parallel processing (MPP)
- D. streaming

Correct Answer: D

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-in/azure/azure-monitor/overview>

QUESTION 71

You have a SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours.

You need to minimize the cost of the SQL pool during idle times. The solution must ensure that the data remains intact.

What should you do on the SQL pool?

- A. Scale down the data warehouse units (DWUs).
- B. Pause the pool.
- C. Create a user-defined restore point.
- D. Delete the pool

Correct Answer: B

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-compute-overview>

QUESTION 72

Which Azure Data Factory component initiates the execution of a pipeline?

- A. a control flow
- B. a trigger
- C. a parameter
- D. an activity

Correct Answer: B

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipeline-execution-triggers#trigger-execution>

QUESTION 73

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse.

Which type of processing does the reporting solution use?

- A. stream processing
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. Online Transaction Processing (OLTP)

Correct Answer: C

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://datawarehouseinfo.com/how-does-oltp-differ-from-olap-database/>

QUESTION 74

DRAG DROP

Match the types of activities to the appropriate Azure Data Factory activities.

To answer, drag the appropriate activity type from the column on the left to its Data Factory activity on the right. Each activity type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Activity Types

Control
Data movement
Data transformation

Answer Area

	Copy
	Mapping data flow
	Until

Correct Answer:

Activity Types

Control
Data movement
Data transformation

Answer Area

Data movement	Copy
Data transformation	Mapping data flow
Control	Until

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipelines-activities#data-transformation-activities>

QUESTION 75

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. denormalized data
- B. heavy writes and moderate reads
- C. light writes and heavy reads
- D. schema on write
- E. schema on read
- F. normalized data

Correct Answer: BDF

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Explanation:

B: Transactional data tends to be heavy writes, moderate reads.

D: Typical traits of transactional data include: schema on write, strongly enforced

F: Transactional data tends to be highly normalized.

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

QUESTION 76

Which two activities can be performed entirely by using the Microsoft Power BI service? Each correct answer presents a complete solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. report and dashboard creation
- B. report sharing and distribution
- C. data modeling
- D. data acquisition and preparation

Correct Answer: AD

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

QUESTION 77

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

In Azure Data Factory, you can use

- | |
|------------------------|
| a control flow |
| a dataset |
| a linked service |
| an integration runtime |

activities that depend on the output of other pipeline activities.

Correct Answer:

Answer Area

In Azure Data Factory, you can use

- | |
|------------------------|
| a control flow |
| a dataset |
| a linked service |
| an integration runtime |

activities that depend on the output of other pipeline activities.

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/frequently-asked-questions>

QUESTION 78

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down.

What should you use in Microsoft Power BI?

- A. a dashboard
- B. a report
- C. a dataflow
- D. Microsoft Power Apps

Correct Answer: B

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards>

<https://docs.microsoft.com/en-us/power-bi/visuals/power-bi-visualization-kpi>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-drill>

QUESTION 79

Which two Azure services can be used to provision Apache Spark clusters? Each correct answer presents a complete solution. (Choose two.)

NOTE: Each correct selection is worth one point.

- A. Azure Time Series Insights
- B. Azure HDInsight
- C. Azure Databricks
- D. Azure Log Analytics

Correct Answer: BC

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://www.sqlshack.com/a-beginners-guide-to-azure-databricks/>

QUESTION 80

You have a quality assurance application that reads data from a data warehouse.

Which type of processing does the application use?

- A. Online Transaction Processing (OLTP)
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. stream processing

Correct Answer: A

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>**QUESTION 81**

Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution. (Choose three.)

NOTE: Each correct selection is worth one point.

- A. a report page
- B. a Microsoft PowerPoint slide
- C. a visualization from a report
- D. a dataflow
- E. a text box

Correct Answer: ACE

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards><https://docs.microsoft.com/en-us/power-bi/create-reports/service-dashboard-add-widget>**QUESTION 82**

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/power-bi/connect-data/service-datasets-across-workspaces>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards><https://powerbi.microsoft.com/en-us/excel-and-power-bi/>

QUESTION 83

Which Azure Data Factory component provides the compute environment for activities?

- A. a linked service
- B. an integration runtime
- C. a control flow
- D. a pipeline

Correct Answer: B

Section: Describe an analytics workload on Azure
Explanation

Explanation/Reference:

Explanation:

The Integration Runtime (IR) is the compute infrastructure used by Azure Data Factory to provide the following data integration capabilities across different network environments:

- Data Flow
- Data movement
- Activity dispatch
- SSIS package execution

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-integration-runtime>

QUESTION 84

You need to use Transact-SQL to query files in Azure Data Lake Storage from an Azure Synapse Analytics data warehouse.

What should you use to query the files?

- A. Azure Functions
- B. Microsoft SQL Server Integration Services (SSIS)
- C. PolyBase
- D. Azure Data Factory

Correct Answer: C

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.databricks.com/data/data-sources/azure/synapse-analytics.html>

QUESTION 85

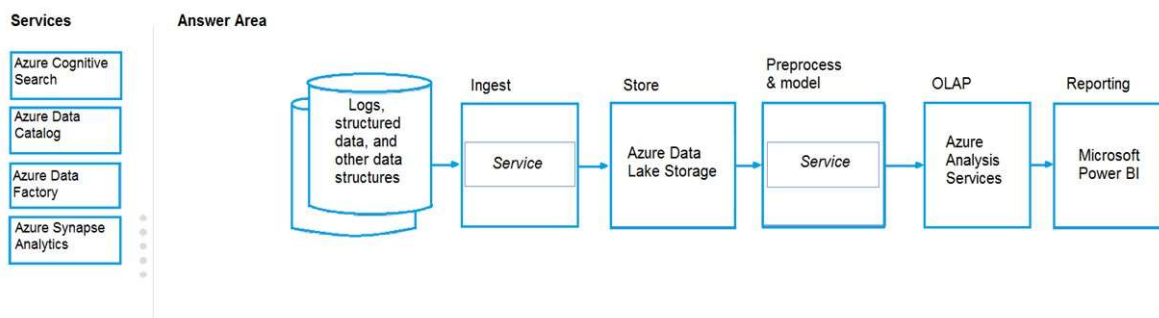
DRAG DROP

Match the Azure services to the appropriate locations in the architecture.

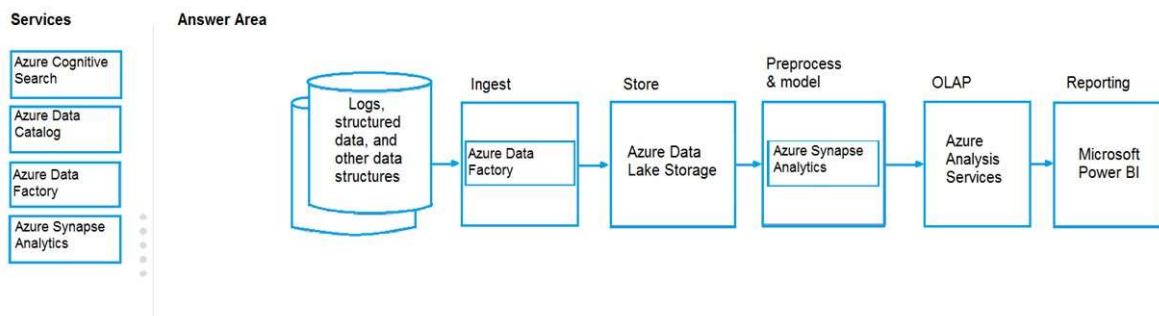
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:



Correct Answer:



Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Box Ingest: Azure Data Factory

You can build a data ingestion pipeline with Azure Data Factory (ADF).

Box Preprocess & model: Azure Synapse Analytics
Use Azure Synapse Analytics to preprocess data and deploy machine learning models.

Reference:

<https://docs.microsoft.com/en-us/azure/machine-learning/how-to-data-ingest-adf>

<https://docs.microsoft.com/en-us/azure/machine-learning/team-data-science-process/sqldw-walkthrough>

QUESTION 86

DRAG DROP

Match the types of workloads to the appropriate scenarios.

To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Workload Types

Batch
Streaming

Answer Area

Workload type	Analyzing web traffic from the past year
Workload type	Classifying images that were uploaded last month
Workload type	Tracking how many people are currently using a website

Correct Answer:

Workload Types

Batch
Streaming

Answer Area

Batch	Analyzing web traffic from the past year
Batch	Classifying images that were uploaded last month
Streaming	Tracking how many people are currently using a website

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: Batch

The batch processing model requires a set of data that is collected over time while the stream processing model requires data to be fed into an analytics tool, often in micro-batches, and in real-time.

The batch Processing model handles a large batch of data while the Stream processing model handles individual records or micro-batches of few records.

In Batch Processing, it processes over all or most of the data but in Stream Processing, it processes over data

on a rolling window or most recent record.

Box 2: Batch

Box 3: Streaming

Reference:

<https://k21academy.com/microsoft-azure/dp-200/batch-processing-vs-stream-processing>

QUESTION 87

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

In a data warehousing workload, data

	▼
from a single source is distributed to multiple locations	
from multiple sources is combined in a single location	
is added to a queue for multiple systems to process	
is used to train machine learning models	

Correct Answer:

Answer Area

In a data warehousing workload, data

	▼
from a single source is distributed to multiple locations	
from multiple sources is combined in a single location	
is added to a queue for multiple systems to process	
is used to train machine learning models	

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Note: The data warehouse workload encompasses:

- The entire process of loading data into the warehouse
- Performing data warehouse analysis and reporting
- Managing data in the data warehouse
- Exporting data from the data warehouse

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-management>

QUESTION 88

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	<input type="radio"/>	<input type="radio"/>
Azure Data Factory pipelines can execute other pipelines	<input type="radio"/>	<input type="radio"/>
A processing step within an Azure Data Factory pipeline is an activity	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	<input type="radio"/>	<input checked="" type="radio"/>
Azure Data Factory pipelines can execute other pipelines	<input checked="" type="radio"/>	<input type="radio"/>
A processing step within an Azure Data Factory pipeline is an activity	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: No

A pipeline is a logical grouping of activities that together perform a task.

Box 2: Yes

You can construct pipeline hierarchies with data factory.

Box 3: Yes

A pipeline is a logical grouping of activities that together perform a task.

Reference:

<https://mrpaulandrew.com/2019/09/25/azure-data-factory-pipeline-hierarchies-generation-control/>

QUESTION 89

DRAG DROP

Match the Azure services to the appropriate requirements.

To answer, drag the appropriate service from the column on the left to its requirement on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:

Services

Azure Data Factory

Azure Data Lake Storage

Azure SQL Database

Azure Synapse Analytics

Answer Area

Service

Output data to Parquet format

Service

Store data that is in Parquet format

Service

Persist a tabular representation of data that is stored in Parquet format

Correct Answer:

Services

Azure Data Factory

Azure Data Lake Storage

Azure SQL Database

Azure Synapse Analytics

Answer Area

Azure Data Factory

Output data to Parquet format

Azure Data Lake Storage

Store data that is in Parquet format

Azure Synapse Analytics

Persist a tabular representation of data that is stored in Parquet format

Section: Describe an analytics workload on Azure Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Data Factory

Box 2: Azure Data Lake Storage

Azure Data Lake Storage (ADLA) now natively supports Parquet files. ADLA adds a public preview of the native extractor and outputter for the popular Parquet file format

Box 3: Azure Synapse Analytics

Use Azure Synapse Analytics Workspaces.

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/supported-file-formats-and-compression-codecs>

QUESTION 90

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	<input type="radio"/>	<input type="radio"/>
Azure Synapse Analytics can be paused to reduce compute costs	<input type="radio"/>	<input type="radio"/>
An Azure Synapse Analytics data warehouse has a fixed storage capacity	<input type="radio"/>	<input type="radio"/>

Correct Answer:

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	<input checked="" type="radio"/>	<input type="radio"/>
Azure Synapse Analytics can be paused to reduce compute costs	<input checked="" type="radio"/>	<input type="radio"/>
An Azure Synapse Analytics data warehouse has a fixed storage capacity	<input type="radio"/>	<input checked="" type="radio"/>

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: Yes

Compute is separate from storage, which enables you to scale compute independently of the data in your system.

Box 2: Yes

You can use the Azure portal to pause and resume the dedicated SQL pool compute resources. Pausing the data warehouse pauses compute. If your data warehouse was paused for the entire hour, you will not be charged compute during that hour.

Box 3: No

Storage is sold in 1 TB allocations. If you grow beyond 1 TB of storage, your storage account will automatically grow to 2 TBs.

Reference:

<https://azure.microsoft.com/en-us/pricing/details/synapse-analytics/>

QUESTION 91

What should you use to build a Microsoft Power BI paginated report?

- A. Charticulator
- B. Power BI Desktop
- C. the Power BI service
- D. Power BI Report Builder

Correct Answer: D

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Power BI Report Builder is the standalone tool for authoring paginated reports for the Power BI service.

Reference:

<https://docs.microsoft.com/en-us/power-bi/paginated-reports/paginated-reports-report-builder-power-bi>

QUESTION 92

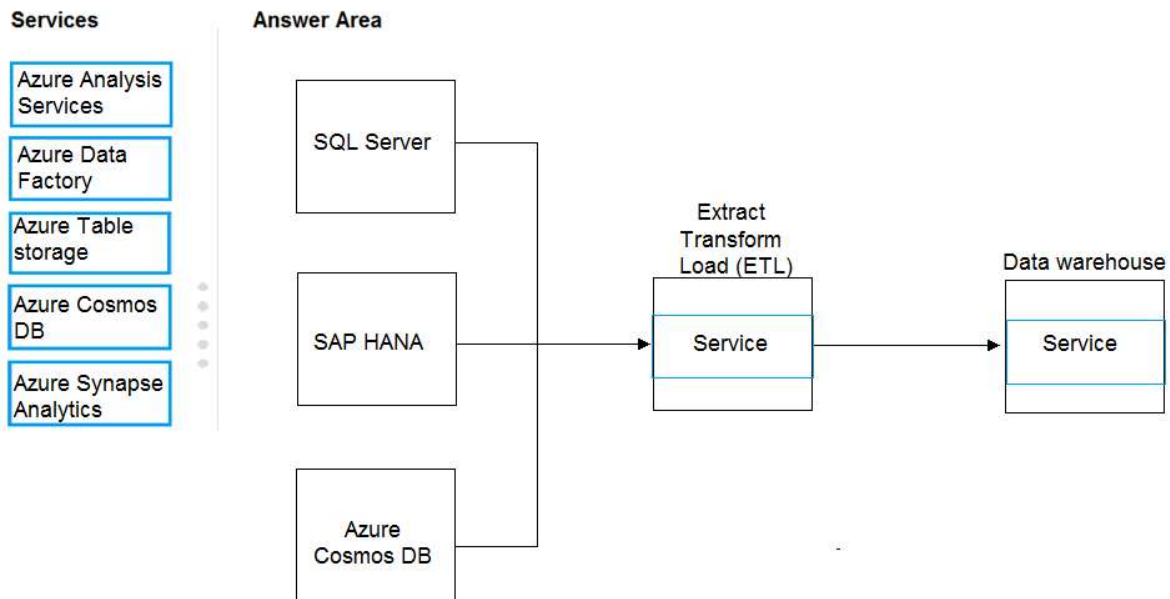
DRAG DROP

Match the Azure services to the appropriate locations in the architecture.

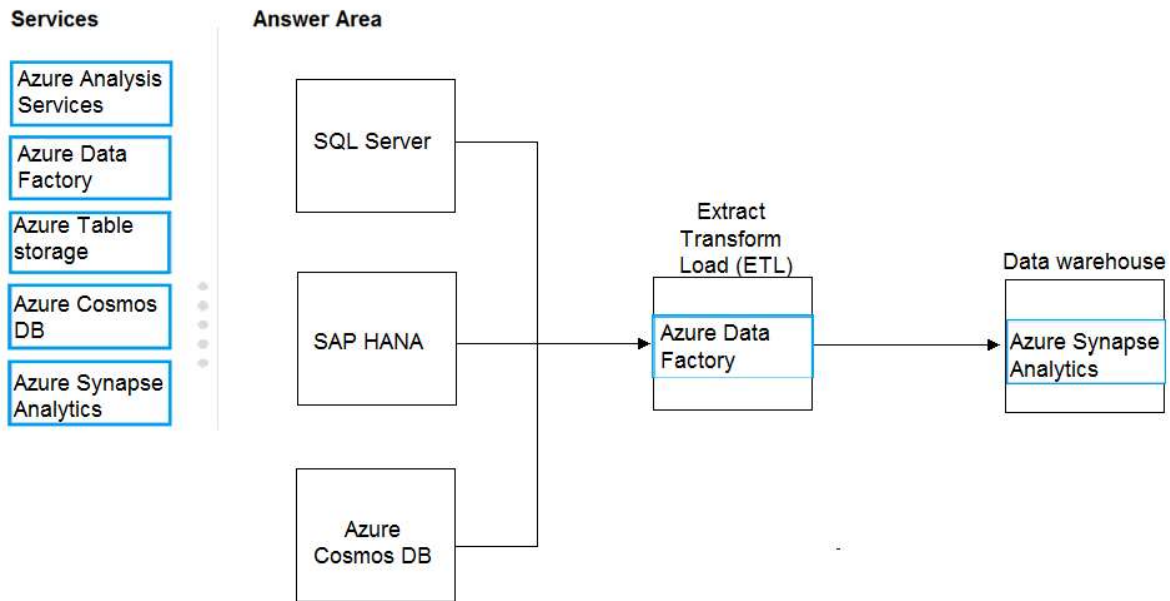
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.

NOTE: Each correct match is worth one point.

Select and Place:



Correct Answer:



Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

Explanation:

Box 1: Azure Data factory

Relevant Azure service for the three ETL phases are Azure Data Factory and SQL Server Integration Services (SSIS).

Box 2: Azure Synapse Analytics

You can copy and transform data in Azure Synapse Analytics by using Azure Data Factory

Note: Azure Synapse Analytics connector is supported for the following activities:

- Copy activity with supported source/sink matrix table
- Mapping data flow
- Lookup activity
- GetMetadata activity

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl>

<https://docs.microsoft.com/en-us/azure/data-factory/connector-azure-sql-data-warehouse>